

Essential Mirror — Archetype Specification

INTERNAL. Not published, not linked, not in frontend/public/.

Derived from `backend/services/essential_scoring.py` and `backend/constants/essential_data.json` as shipped, June 2026. The reader-facing companion is `docs/ARCHETYPES_GUIDE.md`, which deliberately omits sections 4–7 of this document.

Nothing in this document changed any score. It records what the code does, including three defects found while writing it.

1. Structure

Structure	Value
Archetypes	6
Items per lens	50
Lenses	2 — <code>self_assessment</code> , <code>ideal_partner</code>
Items scored per archetype	10
Response scale	1–5 Likert
Output per archetype	raw score, and percentage of that archetype's own maximum
Primary / secondary	highest and second-highest percentage within a lens
Evidence tier	developmental

Key-to-name mapping is not one-to-one with the display names, which matters when reading logs:

Internal key	Display name	Subtitle	Public slug
<code>rock</code>	The Rock	Stability & Tradition	<code>the-rock</code>
<code>challenger</code>	The Challenger	Intellect & Growth	<code>the-challenger</code>
<code>empath</code>	The Empath	Connection & Intimacy	<code>the-empath</code>
<code>adventurer</code>	The Voyager	Social & Lifestyle	<code>the-voyager</code>
<code>diplomat</code>	The Diplomat	Harmony & Communication	<code>the-diplomat</code>
<code>visionary</code>	The Torchbearer	Purpose & Ambition	<code>the-torchbearer</code>

2. Item allocation as shipped

Archetype	Items	Reverse-scored
rock	3, 4, 6, 12, 13, 15, 29, 45, 47, 49	—
challenger	1, 10, 24, 31, 32, 33, 35, 39, 40, 50	—
empath	8, 21, 27, 28, 41, 42, 43, 44, 46, 48	48
adventurer	2, 5, 9, 11, 14, 16, 17, 18, 19, 20	11, 16, 20
diplomat	17, 22, 23, 25, 26, 27, 28, 29, 30, 38	24 (<i>inert</i> — see §6)
visionary	1, 2, 5, 9, 10, 33, 35, 37, 40, 44	—

Reverse scoring is $6 - \text{raw}$, applied within the archetype whose key lists the item.

Per-archetype maximum is $10 \times 5 = 50$; the percentage is $\text{round}(100 \times \text{raw} / 50)$.

3. Shadow and counter-type map

Archetype	Shadow	Counter-type
rock	empath	The Free Spirit
challenger	diplomat	The Harmonizer
empath	challenger	The Stoic
adventurer	rock	The Homebody
diplomat	visionary	The Provocateur
visionary	rock	The Drifter

Properties of the map, which are asymmetric by design or by accident — undetermined:

- `empath` ↔ `challenger` is the only **mutual** pull.
- `rock` is the shadow of **two** archetypes (`adventurer`, `visionary`).
- `adventurer` is **nobody's** shadow. No reader is ever told the Voyager is their shadow pull.

The third point is worth a decision. If the map is meant to be a complete function over the set, one archetype being unreachable is a gap; if it is meant to be an editorial judgement per archetype, it is fine and should be recorded as intentional.

4. DEFECT — item overlap between archetypes

Six archetypes × 10 items = **60 scoring slots drawn from a 50-item lens**. Thirteen items are therefore counted toward two archetypes each. The archetype scores are not independent, and no part of the shipped output says so.

Pairwise overlap:

Pair	Shared items	Items
challenger × visionary	5	1, 10, 33, 35, 40
adventurer × visionary	3	2, 5, 9
empath × diplomat	2	27, 28
rock × diplomat	1	29
adventurer × diplomat	1	17
empath × visionary	1	44

The Torchbearer is the serious case. Eight of its ten items are shared — five with the Challenger and three with the Voyager. Only items 37 and 44 sit outside that, and 44 is shared with the Empath. In practice `visionary` is very largely a linear combination of `challenger` and `adventurer`, so:

- a reader who scores high on Challenger is *mechanically* pushed up on Torchbearer;
- "primary: Torchbearer, secondary: Challenger" is a near-guaranteed pairing rather than a finding;
- the Delta on Torchbearer partly restates the Delta on Challenger.

The Diplomat is the second case. Four of its ten items are shared, across three other archetypes.

Consequences to hold in mind:

1. Any claim that the six archetypes are distinct patterns is not currently supported by the item allocation.
2. Primary-archetype assignment can be driven by shared items rather than by distinguishing ones.
3. This interacts with the MRD work: gating differences between correlated archetype scores is a different and harder problem than gating differences between independent ones.

Deliberately not fixed. Re-allocating items changes scores, which is a frozen-scoring decision. Under the immutability rule, any re-allocation ships as a new scoring version and every existing report keeps its original numbers. Decision deferred; overlap documented internally only, per instruction.

5. DEFECT — three items score nothing

Items **7, 34 and 36** appear in no archetype's item list. A reader answers them twice, once per lens, and they contribute to no archetype score.

They are not necessarily wasted: `calculate_dimension_scores` computes five display dimensions over its own item groupings, and these items may fall there. But as far as the six archetypes are concerned, six of a reader's hundred answers do nothing.

Either they should be allocated, or they should be identified as dimension-only items so it is clear the omission is intentional.

6. DEFECT — the Diplomat's reverse key is inert

`diplomat.reverse = [24]`, but item 24 is not in `diplomat.questions`. Reverse scoring is applied only to items within the archetype's own list, so this entry never fires.

Item 24 is in `challenger.questions`, where `reverse` is empty. So one of two things is true and we do not currently know which:

- item 24 was intended to be reverse-scored for the Challenger and the entry landed on the wrong archetype; or
- item 24 was intended to be a Diplomat item and was never added to the list.

Both readings imply the shipped Challenger score is missing an intended reversal. **No fix applied** — it would change scores. Recorded for the same versioned decision as §4.

A guard is cheap and worth adding regardless of the outcome: assert at import that every entry in `reverse` appears in the same archetype's `questions`, so this class of defect cannot ship again.

7. Naming risk — the `compatibility` field

`get_compatibility_result()` returns a **blend narrative for the top two archetypes within a single lens**. It contains no number, no second person, and no comparison between people, so it does not breach the permanent no-compatibility-score refusal.

But the field is named `compatibility` in the API response, in the stored result, and in the PDF builder, and the lookup table is `_COMPATIBILITY`. Anyone reading the payload — a partner, a practitioner, a journalist, a future engineer — will reasonably read that as the thing we have said we will never produce.

Recommend renaming to `blend`, carrying the old field in parallel for existing snapshots. Raised with the user; no instruction received; **not changed**.

8. Open questions

1. Re-allocate the 50 items for archetype independence? Requires a scoring version bump, parity tests, and old reports keeping their numbers. (§4)
2. Allocate items 7, 34, 36, or declare them dimension-only? (§5)
3. Resolve item 24's intended reversal. (§6)
4. Rename the `compatibility` field. (§7)
5. Is the shadow map meant to be complete? `adventurer` is currently unreachable as a shadow. (§3)
6. Should the reader-facing guide disclose the overlap? Currently **no**, by instruction, pending the decision in §4. If the overlap is not fixed, this should be revisited — the evidence-tier disclosures set a standard that this omission sits awkwardly against.